

SDS 384 11: Theoretical Statistics

Lecture 1: Introduction

Purnamrita Sarkar
Department of Statistics and Data Science
The University of Texas at Austin
<https://psarkar.github.io/teaching>

Manegerial Stuff

- Instructor- Purnamrita Sarkar
- Course material and homeworks will be posted under www.cs.cmu.edu/~psarkar/teaching/sds384.html
- Office hours: Tuesdays 2-3 pm. GDC 7.504
- TA: Jack Liu
- Homeworks are due Biweekly after class on Wednesdays
- Grading - 5 homeworks (70%), class participation (5%) Final Exam (25%)
- Books
 - Asymptotic Statistics, Aad van der Vaart. Cambridge. 1998.
 - Martin Wainwright's High dimensional statistics: A non-asymptotic view point

Why do theory?

- Say you have estimated $\hat{\theta}_n$ from data X_1, \dots, X_n . How do we know we have a “good” estimation method?
 - Does $\hat{\theta}_n \rightarrow \theta$? This brings us to **Stochastic Convergence**.
- How do I know if one estimation method is better than another?
 - Does the estimate from one converge faster than the other?
 - Does one algorithm work under broader parameter regimes, or weaker assumptions?
 - What is the optimal rate for a given estimation problem?

This class

Your instructor “hopes to cover”:

- Consistency of parameter estimates
 - Stochastic Convergence
 - Concentration inequalities
 - Asymptotic normality of estimators
- Empirical processes, VC classes, covering numbers
- Asymptotic testing
- Examples of network clustering with a bit of random matrix theory
- Bootstrap, Nonparametric regression and density estimation

Stochastic Convergence

Assume that $X_n, n \geq 1$ and X are elements of a separable metric space (S, d) .

Definition (Weak Convergence)

A sequence of random variables converge in “law” or in “distribution” to a random variable X , i.e. $X_n \xrightarrow{d} X$ if $P(X_n \leq x) \rightarrow P(X \leq x) \forall x$ at which $P(X \leq x)$ is continuous.

Definition (Convergence in Probability)

A sequence of random variables converge in “probability” to a random variable X , i.e. $X_n \xrightarrow{P} X$ if $\forall \epsilon > 0, P(d(X_n, X) \geq \epsilon) \rightarrow 0$.

Stochastic Convergence

Assume that $X_n, n \geq 1$ and X are elements of a separable metric space (S, d) .

Definition (Almost Sure Convergence)

A sequence of random variables converges almost surely to a random variable X , i.e. $X_n \xrightarrow{a.s.} X$ if $P\left(\lim_{n \rightarrow \infty} d(X_n, X) = 0\right) = 1$.

- If you think about a (scalar) random variable as a function that maps events to a real number, almost sure convergence means
$$P(\omega \in \Omega : \lim_{n \rightarrow \infty} X_n(\omega) = X(\omega)) = 1$$

Definition (Convergence in quadratic mean)

A sequence of random variables converges in quadratic mean to a random variable X , i.e. $X_n \xrightarrow{q.m.} X$ if $E\left[d(X_n, X)^2\right] \rightarrow 0$.

Theorem

$$X_n \xrightarrow{a.s.} X, X_n \xrightarrow{q.m.} X \Rightarrow X_n \xrightarrow{P} X \Rightarrow X_n \xrightarrow{d} X$$
$$X_n \xrightarrow{d} c \Rightarrow X_n \xrightarrow{P} c$$

Converses: $X_n \xrightarrow{d} X \not\Rightarrow X_n \xrightarrow{P} X$

- Convergence in law needs no knowledge of the joint distribution of X_n and the limiting random variable X .
- Convergence in probability does.

Example

Consider $X \sim N(0, 1)$, $X_n = -X$. $X_n \xrightarrow{d} X$. But how about $X_n \xrightarrow{P} X$?

Converses: $X_n \xrightarrow{d} X \not\Rightarrow X_n \xrightarrow{P} X$

- Convergence in law needs no knowledge of the joint distribution of X_n and the limiting random variable X .
- Convergence in probability does.

Example

Consider $X \sim N(0, 1)$, $X_n = -X$. $X_n \xrightarrow{d} X$. But how about $X_n \xrightarrow{P} X$?

- $P(|X_n - X| \geq \epsilon) = P(2|X| \geq \epsilon) \not\rightarrow 0 \forall \epsilon > 0$. So X_n does not converge in probability to X .

Example

Example

Let $Z \sim U(0, 1)$ and for $n = 2^k + m$ for $k \geq 0, 0 \leq m < 2^k$
 $X_n = 1(Z \in [m2^{-k}, (m+1)2^{-k}])$, i.e. $X_1 = 1$, $X_2 = 1(Z \in [0, 1/2))$,
 $X_3 = 1(Z \in [1/2, 1))$, $X_4 = 1(Z \in [0, 1/4))$, $X_5 = 1(Z \in [1/4, 1/2))$.

Example

Example

Let $Z \sim U(0, 1)$ and for $n = 2^k + m$ for $k \geq 0, 0 \leq m < 2^k$
 $X_n = 1(Z \in [m2^{-k}, (m+1)2^{-k}])$, i.e. $X_1 = 1$, $X_2 = 1(Z \in [0, 1/2))$,
 $X_3 = 1(Z \in [1/2, 1))$, $X_4 = 1(Z \in [0, 1/4))$, $X_5 = 1(Z \in [1/4, 1/2))$.

- For any $Z \in (0, 1)$, the sequence $\{X_n(Z)\}$ does not converge. So $X_n \not\overset{a.s.}{\rightarrow} 0$.

Example

Example

Let $Z \sim U(0, 1)$ and for $n = 2^k + m$ for $k \geq 0, 0 \leq m < 2^k$
 $X_n = 1(Z \in [m2^{-k}, (m+1)2^{-k}])$, i.e. $X_1 = 1$, $X_2 = 1(Z \in [0, 1/2))$,
 $X_3 = 1(Z \in [1/2, 1))$, $X_4 = 1(Z \in [0, 1/4))$, $X_5 = 1(Z \in [1/4, 1/2))$.

- For any $Z \in (0, 1)$, the sequence $\{X_n(Z)\}$ does not converge. So $X_n \not\stackrel{a.s.}{\rightarrow} 0$.
- X_n are a sequence of bernoulli's with probabilities $p_n = 1/2^k$ where $k = \lfloor \log n \rfloor$.

Example

Example

Let $Z \sim U(0, 1)$ and for $n = 2^k + m$ for $k \geq 0, 0 \leq m < 2^k$
 $X_n = 1(Z \in [m2^{-k}, (m+1)2^{-k}])$, i.e. $X_1 = 1$, $X_2 = 1(Z \in [0, 1/2))$,
 $X_3 = 1(Z \in [1/2, 1))$, $X_4 = 1(Z \in [0, 1/4))$, $X_5 = 1(Z \in [1/4, 1/2))$.

- For any $Z \in (0, 1)$, the sequence $\{X_n(Z)\}$ does not converge. So $X_n \not\stackrel{a.s.}{\rightarrow} 0$.
- X_n are a sequence of bernoulli's with probabilities $p_n = 1/2^k$ where $k = \lfloor \log n \rfloor$.
- So $X_n \xrightarrow{P} 0$ and $X_n \xrightarrow{qm} 0$

Example

Example

Let $Z \sim U([0, 1])$ and $X_n = 2^n 1(Z \in [0, 1/n])$. Does X_n converge to X in quadratic mean, almost surely or in probability?

Example

Example

Let $Z \sim U([0, 1])$ and $X_n = 2^n 1(Z \in [0, 1/n])$. Does X_n converge to X in quadratic mean, almost surely or in probability?

- $P(\lim_{n \rightarrow \infty} X_n = X) = P(Z > 0) = 1$. So $X_n \xrightarrow{a.s.} X$.

Example

Example

Let $Z \sim U([0, 1])$ and $X_n = 2^n 1(Z \in [0, 1/n])$. Does X_n converge to X in quadratic mean, almost surely or in probability?

- $P(\lim_{n \rightarrow \infty} X_n = X) = P(Z > 0) = 1$. So $X_n \xrightarrow{a.s.} X$.
- $E|X_n|^2 = 2^{2n}/n \rightarrow \infty$. So $X_n \not\xrightarrow{qm} 0$

Example

Example

Let $Z \sim U([0, 1])$ and $X_n = 2^n 1(Z \in [0, 1/n])$. Does X_n converge to X in quadratic mean, almost surely or in probability?

- $P(\lim_{n \rightarrow \infty} X_n = X) = P(Z > 0) = 1$. So $X_n \xrightarrow{a.s.} X$.
- $E|X_n|^2 = 2^{2n}/n \rightarrow \infty$. So $X_n \not\xrightarrow{qm} 0$
- $P(|X_n| \geq \epsilon) = P(X_n = 2^n) = P(Z \in [0, 1/n]) = 1/n \rightarrow 0$

Continuous Mapping Theorem

Theorem

Let g be continuous on a set C where $P(X \in C) = 1$. Then,

$$X_n \xrightarrow{d} X \Rightarrow g(X_n) \xrightarrow{d} g(X)$$

$$X_n \xrightarrow{P} X \Rightarrow g(X_n) \xrightarrow{P} g(X)$$

$$X_n \xrightarrow{a.s.} X \Rightarrow g(X_n) \xrightarrow{a.s.} g(X)$$

Example

Let $X_n \xrightarrow{d} X$ where $X \sim N(0, 1)$. Then $X_n^2 \xrightarrow{d} ?$

Example

Let $X_n \xrightarrow{d} X$ where $X \sim N(0, 1)$. Then $X_n^2 \xrightarrow{d} ?$

- Use $g(x) = x^2$.

Example

Let $X_n \xrightarrow{d} X$ where $X \sim N(0, 1)$. Then $X_n^2 \xrightarrow{d} ?$

- Use $g(x) = x^2$.
- Use $X^2 \sim \chi_1^2$.

Example

Let $X_n \xrightarrow{d} X$ where $X \sim N(0, 1)$. Then $X_n^2 \xrightarrow{d} ?$

- Use $g(x) = x^2$.
- Use $X^2 \sim \chi_1^2$.
- So $X_n^2 \xrightarrow{d} \chi_1^2$

Example-continuity points

Let X_1, \dots, X_n be i.i.d. with mean μ and variance σ^2 . We have $\bar{X}_n - \mu \xrightarrow{d} 0$. Consider $g(x) = 1_{x>0}$. Then $g((\bar{X}_n - \mu)^2) \xrightarrow{d} ?$

Example-continuity points

Let X_1, \dots, X_n be i.i.d. with mean μ and variance σ^2 . We have $\bar{X}_n - \mu \xrightarrow{d} 0$. Consider $g(x) = 1_{x>0}$. Then $g((\bar{X}_n - \mu)^2) \xrightarrow{d} ?$

- Using Continuous Mapping Theorem, $(\bar{X}_n - \mu)^2 \xrightarrow{d} 0$

Example-continuity points

Let X_1, \dots, X_n be i.i.d. with mean μ and variance σ^2 . We have $\bar{X}_n - \mu \xrightarrow{d} 0$. Consider $g(x) = 1_{x>0}$. Then $g((\bar{X}_n - \mu)^2) \xrightarrow{d} ?$

- Using Continuous Mapping Theorem, $(\bar{X}_n - \mu)^2 \xrightarrow{d} 0$
- Can we use Continuous Mapping Theorem to claim that $g(\bar{X}_n - \mu)^2 \xrightarrow{d} 0$?

Example-continuity points

Let X_1, \dots, X_n be i.i.d. with mean μ and variance σ^2 . We have $\bar{X}_n - \mu \xrightarrow{d} 0$. Consider $g(x) = 1_{x>0}$. Then $g((\bar{X}_n - \mu)^2) \xrightarrow{d} ?$

- Using Continuous Mapping Theorem, $(\bar{X}_n - \mu)^2 \xrightarrow{d} 0$
- Can we use Continuous Mapping Theorem to claim that $g(\bar{X}_n - \mu)^2 \xrightarrow{d} 0$?
- NO. Because, 0 is a random variable whose mass is at 0, where g is discontinuous.

How about convergence in q.m.?

- If $X_n \xrightarrow{qm} X$, then is it true that for continuous f (discontinuous only at a measure zero set), $f(X_n) \xrightarrow{qm} f(X)$?

How about convergence in q.m.?

- If $X_n \xrightarrow{qm} X$, then is it true that for continuous f (discontinuous only at a measure zero set), $f(X_n) \xrightarrow{qm} f(X)$?
- Consider an L -Lipschitz function $f(X)$. $|f(x) - f(y)| \leq L|x - y|$.

How about convergence in q.m.?

- If $X_n \xrightarrow{qm} X$, then is it true that for continuous f (discontinuous only at a measure zero set), $f(X_n) \xrightarrow{qm} f(X)$?
- Consider an L - Lipschitz function $f(X)$. $|f(x) - f(y)| \leq L|x - y|$.
- $E[|f(X_n) - f(X)|^2] \leq L^2 E[|X_n - X|^2] \rightarrow 0$. So for Lipschitz functions quadratic mean convergence goes through.

How about convergence in q.m.?

- If $X_n \xrightarrow{qm} X$, then is it true that for continuous f (discontinuous only at a measure zero set), $f(X_n) \xrightarrow{qm} f(X)$?
- Consider an L -Lipschitz function $f(X)$. $|f(x) - f(y)| \leq L|x - y|$.
- $E[|f(X_n) - f(X)|^2] \leq L^2 E[|X_n - X|^2] \rightarrow 0$. So for Lipschitz functions quadratic mean convergence goes through.
- Can you come up with a non-Lipschitz function and a sequence $\{X_n\}$ where $f(X_n) \not\xrightarrow{qm} 0$?

Portmanteau Theorem

Theorem

The following are equivalent.

- $X_n \xrightarrow{d} X$
- $E[f(X_n)] \rightarrow E[f(X)]$ for all bounded and continuous f .
- $E[f(X_n)] \rightarrow E[f(X)]$ for all bounded and Lipschitz f .
- $E[e^{it^T X_n}] \rightarrow E[e^{it^T X}]$, $\forall t \in \mathbb{R}^k$. (Levy's continuity theorem)
- $t^T X_n \xrightarrow{d} t^T X \quad \forall t \in \mathbb{R}^k$. (Cramer-Wold device)
- $\liminf_n E[f(X_n)] \geq E[f(X)]$ for all non-negative continuous f
- $\limsup_n P(X_n \in F) \leq P(X \in F)$ for all closed F
- $\liminf_n P(X_n \in F) \geq P(X \in F)$ for all open F
- $P(X_n \in B) \rightarrow P(X \in B)$ for all continuity sets B ($P(X \in \partial B) = 0$)

Example-bounded

Consider $f(x) = x$ and

$$X_n = \begin{cases} n & \text{w.p. } 1/n \\ 0 & \text{w.p. } 1 - 1/n \end{cases}$$

Example-bounded

Consider $f(x) = x$ and

$$X_n = \begin{cases} n & \text{w.p. } 1/n \\ 0 & \text{w.p. } 1 - 1/n \end{cases}$$

- $X_n \xrightarrow{d} 0$, but $E[X_n] \rightarrow ?$

Example-bounded

Consider $f(x) = x$ and

$$X_n = \begin{cases} n & \text{w.p. } 1/n \\ 0 & \text{w.p. } 1 - 1/n \end{cases}$$

- $X_n \xrightarrow{d} 0$, but $E[X_n] \rightarrow ?$
- $E[X_n] = 1$. What went wrong?

Example-bounded

Consider $f(x) = x$ and

$$X_n = \begin{cases} n & \text{w.p. } 1/n \\ 0 & \text{w.p. } 1 - 1/n \end{cases}$$

- $X_n \xrightarrow{d} 0$, but $E[X_n] \rightarrow ?$
- $E[X_n] = 1$. What went wrong?
- $f(x) = x$ is not bounded.

Putting everything together

Theorem

$$X_n \xrightarrow{d} X \text{ and } d(X_n, Y_n) \xrightarrow{P} 0 \Rightarrow Y_n \xrightarrow{d} X \quad (1)$$

$$X_n \xrightarrow{d} X \text{ and } Y_n \xrightarrow{d} c \Rightarrow (X_n, Y_n) \xrightarrow{d} (X, c) \quad (2)$$

$$X_n \xrightarrow{P} X \text{ and } Y_n \xrightarrow{P} Y \Rightarrow (X_n, Y_n) \xrightarrow{P} (X, Y) \quad (3)$$

Putting everything together

Theorem

$$X_n \xrightarrow{d} X \text{ and } d(X_n, Y_n) \xrightarrow{P} 0 \Rightarrow Y_n \xrightarrow{d} X \quad (1)$$

$$X_n \xrightarrow{d} X \text{ and } Y_n \xrightarrow{d} c \Rightarrow (X_n, Y_n) \xrightarrow{d} (X, c) \quad (2)$$

$$X_n \xrightarrow{P} X \text{ and } Y_n \xrightarrow{P} Y \Rightarrow (X_n, Y_n) \xrightarrow{P} (X, Y) \quad (3)$$

- Eq 3 does not hold if we replace convergence in probability by convergence in distribution.

Putting everything together

Theorem

$$X_n \xrightarrow{d} X \text{ and } d(X_n, Y_n) \xrightarrow{P} 0 \Rightarrow Y_n \xrightarrow{d} X \quad (1)$$

$$X_n \xrightarrow{d} X \text{ and } Y_n \xrightarrow{d} c \Rightarrow (X_n, Y_n) \xrightarrow{d} (X, c) \quad (2)$$

$$X_n \xrightarrow{P} X \text{ and } Y_n \xrightarrow{P} Y \Rightarrow (X_n, Y_n) \xrightarrow{P} (X, Y) \quad (3)$$

- Eq 3 does not hold if we replace convergence in probability by convergence in distribution.
- Example: $X_n \sim N(0, 1)$, $Y_n = -X_n$. $X \perp Y$ and X, Y are independent standard normal random variables.

Putting everything together

Theorem

$$X_n \xrightarrow{d} X \text{ and } d(X_n, Y_n) \xrightarrow{P} 0 \Rightarrow Y_n \xrightarrow{d} X \quad (1)$$

$$X_n \xrightarrow{d} X \text{ and } Y_n \xrightarrow{d} c \Rightarrow (X_n, Y_n) \xrightarrow{d} (X, c) \quad (2)$$

$$X_n \xrightarrow{P} X \text{ and } Y_n \xrightarrow{P} Y \Rightarrow (X_n, Y_n) \xrightarrow{P} (X, Y) \quad (3)$$

- Eq 3 does not hold if we replace convergence in probability by convergence in distribution.
- Example: $X_n \sim N(0, 1)$, $Y_n = -X_n$. $X \perp Y$ and X, Y are independent standard normal random variables.
- Then $X_n \xrightarrow{d} X$ and $Y_n \xrightarrow{d} Y$. But $(X_n, Y_n) \xrightarrow{d} (X, -X)$, not $(X_n, Y_n) \xrightarrow{d} (X, Y)$.

Putting everything together

Theorem (Slutsky's theorem)

$X_n \xrightarrow{d} X$ and $Y_n \xrightarrow{d} c$ imply that

$$X_n + Y_n \xrightarrow{d} X + c$$

$$X_n Y_n \xrightarrow{d} cX$$

$$X_n / Y_n \xrightarrow{d} X / c$$

Putting everything together

Theorem (Slutsky's theorem)

$X_n \xrightarrow{d} X$ and $Y_n \xrightarrow{d} c$ imply that

$$X_n + Y_n \xrightarrow{d} X + c$$

$$X_n Y_n \xrightarrow{d} cX$$

$$X_n / Y_n \xrightarrow{d} X / c$$

- Does $X_n \xrightarrow{d} X$ and $Y_n \xrightarrow{d} Y$ imply $X_n + Y_n \xrightarrow{d} X + Y$?

Putting everything together

Theorem (Slutsky's theorem)

$X_n \xrightarrow{d} X$ and $Y_n \xrightarrow{d} c$ imply that

$$X_n + Y_n \xrightarrow{d} X + c$$

$$X_n Y_n \xrightarrow{d} cX$$

$$X_n / Y_n \xrightarrow{d} X / c$$

- Does $X_n \xrightarrow{d} X$ and $Y_n \xrightarrow{d} Y$ imply $X_n + Y_n \xrightarrow{d} X + Y$?
- Take $Y_n = -X_n$, and X, Y as independent standard normal random variables. $X_n \xrightarrow{d} X$ and $Y_n \xrightarrow{d} Y$ but $X_n + Y_n \xrightarrow{d} 0$.

Using all this

If X_1, \dots, X_n are i.i.d. random variables with mean μ and variance σ^2 ,
prove that $\sqrt{n} \frac{\bar{X}_n - \mu}{S_n} \xrightarrow{d} N(0, 1)$.

Using all this

If X_1, \dots, X_n are i.i.d. random variables with mean μ and variance σ^2 , prove that $\sqrt{n} \frac{\bar{X}_n - \mu}{S_n} \xrightarrow{d} N(0, 1)$.

- First note that $S_n^2 = \frac{1}{n-1} \sum_i X_i^2 - \bar{X}_n^2 = \frac{n}{n-1} \frac{\sum_i X_i^2}{n} - \bar{X}_n^2$

Using all this

If X_1, \dots, X_n are i.i.d. random variables with mean μ and variance σ^2 , prove that $\sqrt{n} \frac{\bar{X}_n - \mu}{S_n} \xrightarrow{d} N(0, 1)$.

- First note that $S_n^2 = \frac{1}{n-1} \sum_i X_i^2 - \bar{X}_n^2 = \frac{n}{n-1} \frac{\sum_i X_i^2}{n} - \bar{X}_n^2$
- Law of large numbers give $\frac{\sum_i X_i^2}{n} \xrightarrow{P} E[X^2]$ and $\bar{X}_n \xrightarrow{P} \mu$.

Using all this

If X_1, \dots, X_n are i.i.d. random variables with mean μ and variance σ^2 , prove that $\sqrt{n} \frac{\bar{X}_n - \mu}{S_n} \xrightarrow{d} N(0, 1)$.

- First note that $S_n^2 = \frac{1}{n-1} \sum_i X_i^2 - \bar{X}_n^2 = \frac{n}{n-1} \frac{\sum_i X_i^2}{n} - \bar{X}_n^2$
- Law of large numbers give $\frac{\sum_i X_i^2}{n} \xrightarrow{P} E[X^2]$ and $\bar{X}_n \xrightarrow{P} \mu$.
- So $(\frac{\sum_i X_i^2}{n}, \bar{X}_n) \xrightarrow{P} (E[X^2], \mu)$ and now using the continuous mapping theorem, $S_n^2 \xrightarrow{P} \sigma^2$.

Using all this

If X_1, \dots, X_n are i.i.d. random variables with mean μ and variance σ^2 , prove that $\sqrt{n} \frac{\bar{X}_n - \mu}{S_n} \xrightarrow{d} N(0, 1)$.

- First note that $S_n^2 = \frac{1}{n-1} \sum_i X_i^2 - \bar{X}_n^2 = \frac{n}{n-1} \frac{\sum_i X_i^2}{n} - \bar{X}_n^2$
- Law of large numbers give $\frac{\sum_i X_i^2}{n} \xrightarrow{P} E[X^2]$ and $\bar{X}_n \xrightarrow{P} \mu$.
- So $(\frac{\sum_i X_i^2}{n}, \bar{X}_n) \xrightarrow{P} (E[X^2], \mu)$ and now using the continuous mapping theorem, $S_n^2 \xrightarrow{P} \sigma^2$.
- Finally, $\sqrt{n}(\bar{X}_n - \mu) \xrightarrow{d} N(0, \sigma^2)$ using CLT.

Using all this

If X_1, \dots, X_n are i.i.d. random variables with mean μ and variance σ^2 , prove that $\sqrt{n} \frac{\bar{X}_n - \mu}{S_n} \xrightarrow{d} N(0, 1)$.

- First note that $S_n^2 = \frac{1}{n-1} \sum_i X_i^2 - \bar{X}_n^2 = \frac{n}{n-1} \frac{\sum_i X_i^2}{n} - \bar{X}_n^2$
- Law of large numbers give $\frac{\sum_i X_i^2}{n} \xrightarrow{P} E[X^2]$ and $\bar{X}_n \xrightarrow{P} \mu$.
- So $(\frac{\sum_i X_i^2}{n}, \bar{X}_n) \xrightarrow{P} (E[X^2], \mu)$ and now using the continuous mapping theorem, $S_n^2 \xrightarrow{P} \sigma^2$.
- Finally, $\sqrt{n}(\bar{X}_n - \mu) \xrightarrow{d} N(0, \sigma^2)$ using CLT.
- Now using Slutsky's lemma, $\sqrt{n}(\bar{X}_n - \mu)/S_n \xrightarrow{d} N(0, 1)$ using CLT.

Definition

X is defined to be “tight” if $\forall \epsilon > 0 \exists M$ for which,

$$P(\|X\| > M) < \epsilon$$

$\{X_n\}$ is defined to be uniformly tight if $\forall \epsilon > 0 \exists M$ for which,

$$\sup_n P(\|X_n\| > M) < \epsilon$$

Theorem

- $X_n \xrightarrow{d} X \Rightarrow \{X_n\}$ is UT.
- $\{X_n\}$ is UT implies that, there exists a subsequence $\{n_j\}$ such that $X_{n_j} \xrightarrow{d} X$.

Notation for rates, small oh-pee and big oh-pee

Definition

- The small o_P :

$$X_n = o_P(1) \Leftrightarrow X_n \xrightarrow{P} 0$$

$$X_n = o_P(R_n) \Leftrightarrow X_n = Y_n R_n \text{ and } Y_n = o_P(1)$$

X_n is vanishing in probability

- The big O_P :

$$X_n = O_P(1) \Leftrightarrow \{X_n\} \text{ is UT}$$

$$X_n = O_P(R_n) \Leftrightarrow X_n = Y_n R_n \text{ and } Y_n = O_P(1)$$

X_n lies within a ball of finite radius with high probability

How do they interact

$$o_P(a_n) + o_P(b_n) = o_P(\max(a_n, b_n)).$$

$$o_P(a_n) + O_P(b_n) = O_P(b_n).$$

$$O_P(a_n)o_P(b_n) = o_P(a_nb_n).$$

$$1 + O_P(1) = O_P(1).$$

$$(1 + o_P(1))^{-1} = 1 + o_P(1).$$

$$o_P(O_P(1)) = o_P(1).$$

$$X_n = o_P(a_n), |X_n|^r = o_P(a_n^r)$$

Be careful:

$$e^{o_P(1)} \neq o_P(1)$$

$O_P(1) + O_P(1)$ Can actually be $o_P(1)$ because of cancellation.

