

SDS 384 11: Theoretical Statistics

Lecture 19: Overview

Purnamrita Sarkar
Department of Statistics and Data Science
The University of Texas at Austin

Stochastic Convergence

Assume that $X_n, n \geq 1$ and X are elements of a separable metric space (S, d) .

Definition (Weak Convergence)

A sequence of random variable s converge in “law” or in “distribution” to a random variable X , i.e. $X_n \xrightarrow{d} X$ if $P(X_n \leq x) \rightarrow P(X \leq x) \forall x$ at which $P(X \leq x)$ is continuous.

Definition (Convergence in Probability)

A sequence of random variables converge in “probability” to a random variable X , i.e. $X_n \xrightarrow{P} X$ if $\forall \epsilon > 0, P(d(X_n, X) \geq \epsilon) \rightarrow 0$.

Stochastic Convergence

Assume that $X_n, n \geq 1$ and X are elements of a separable metric space (S, d) .

Definition (Almost Sure Convergence)

A sequence of random variables converge almost surely to a random variable X , i.e. $X_n \xrightarrow{a.s.} X$ if $P\left(\lim_{n \rightarrow \infty} d(X_n, X) = 0\right) = 1$.

- If you think about a (scalar) random variable as a function that maps events to a real number, almost sure convergence means $P(\omega \in \Omega : \lim_{n \rightarrow \infty} X_n(\omega) = X(\omega)) = 1$

Definition (Convergence in quadratic mean)

A sequence of random variables converge in quadratic mean to a random variable X , i.e. $X_n \xrightarrow{q.m.} X$ if $E\left[d(X_n, X)^2\right] \rightarrow 0$.

Continuous Mapping Theorem

Theorem

Let g be continuous on a set C where $P(X \in C) = 1$. Then,

$$X_n \xrightarrow{d} X \Rightarrow g(X_n) \xrightarrow{d} g(X)$$

$$X_n \xrightarrow{P} X \Rightarrow g(X_n) \xrightarrow{P} g(X)$$

$$X_n \xrightarrow{a.s.} X \Rightarrow g(X_n) \xrightarrow{a.s.} g(X)$$

- What about continuous mapping with quadratic mean?

Putting everything together

Theorem

$$X_n \xrightarrow{d} X \text{ and } d(X_n, Y_n) \xrightarrow{P} 0 \Rightarrow Y_n \xrightarrow{d} X \quad (1)$$

$$X_n \xrightarrow{d} X \text{ and } Y_n \xrightarrow{d} c \Rightarrow (X_n, Y_n) \xrightarrow{d} (X, c) \quad (2)$$

$$X_n \xrightarrow{P} X \text{ and } Y_n \xrightarrow{P} Y \Rightarrow (X_n, Y_n) \xrightarrow{P} (X, Y) \quad (3)$$

- Eq 3 does not hold if we replace convergence in probability by convergence in distribution.
- Example: $X_n \sim N(0, 1)$, $Y_n = -X_n$. $X \perp Y$ and X, Y are independent standard normal random variables.
- Then $X_n \xrightarrow{d} X$ and $Y_n \xrightarrow{d} Y$. But $(X_n, Y_n) \xrightarrow{d} (X, -X)$, not $(X_n, Y_n) \xrightarrow{d} (X, Y)$.

Putting everything together

Theorem (Slutsky's theorem)

$X_n \xrightarrow{d} X$ and $Y_n \xrightarrow{d} c$ imply that

$$X_n + Y_n \xrightarrow{d} X + c$$

$$X_n Y_n \xrightarrow{d} cX$$

$$X_n / Y_n \xrightarrow{d} X / c$$

- Does $X_n \xrightarrow{d} X$ and $Y_n \xrightarrow{d} Y$ imply $X_n + Y_n \xrightarrow{d} X + Y$?
- Take $Y_n = -X_n$, and X, Y as independent standard normal random variables. $X_n \xrightarrow{d} X$ and $Y_n \xrightarrow{d} Y$ but $X_n + Y_n \xrightarrow{d} 0$.

Lindeberg-feller CLT for triangular arrays

Theorem (Ordinary Central limit theorem)

$X_1, \dots, X_n \stackrel{iid}{\sim} f$ with $E|X_i| \leq \infty$, $E[X_1] = 0$. If $E[X_i^2] = \sigma^2$,
 $\sqrt{n}(\bar{X}_n - \mu) \xrightarrow{d} N(0, \sigma^2)$.

X_{11}

X_{21}, X_{22}

X_{21}, X_{22}, X_{23}

...

Theorem (Lindeberg-feller)

For each n let $(X_{ni})_{i=1}^n$ be independent random variables with mean zero and variance σ_{ni}^2 . Let $Z_n = \sum_{i=1}^n X_{ni}$ and $B_n^2 = \text{var}(Z_n)$. Then

$Z_n/B_n \xrightarrow{d} N(0, 1)$, as long as the **Lindeberg condition** holds.

The Lindeberg condition

Definition (Lindeberg condition)

For every $\epsilon > 0$,

$$\frac{1}{B_n^2} \sum_{j=1}^n E[X_{nj}^2 1(|X_{nj}| \geq \epsilon B_n)] \rightarrow 0 \text{ as } n \rightarrow \infty \quad (4)$$

Converse: If $\frac{\sigma_{nj}^2}{B_n^2} \rightarrow 0$ as $n \rightarrow \infty$, i.e. no one variance plays a significant role in the limit, and if $Z_n/B_n \xrightarrow{d} N(0,1)$, then the Lindeberg condition holds.

Necessary and Sufficient: If $\frac{\sigma_{nj}^2}{B_n^2} \rightarrow 0$, the the Lindeberg condition is necessary and sufficient to show the CLT.

Example

Let X_1, \dots, X_n be independent random variables with mean zero and variance one. Do you think $\sqrt{n}\bar{X}_n \xrightarrow{d} N(0, 1)$?

•

$$X_{nj} = \begin{cases} 2j & \text{w.p. } \frac{1}{8j^2} \\ 0 & \text{w.p. } 1 - \frac{1}{4j^2} \\ -2j & \text{w.p. } \frac{1}{8j^2} \end{cases}$$

- $E[X_{nj}] = 0$ and $\text{var}(X_{nj}) = 1$. $B_n^2 = n$.
- Lets check the Lindeberg condition with $\epsilon = 1$.

$$\frac{1}{n} \sum_j E[X_{nj}^2 1(|X_{nj}| \geq \sqrt{n})] = \frac{1}{n} \sum_j 2 \times 4j^2 1(2j \geq \sqrt{n}) \frac{1}{8j^2} = \frac{1}{n} \sum_{j \geq \sqrt{n}/2} 1 \rightarrow 1$$

- Since $\sigma_{nj}^2/B_n^2 = 1/n \rightarrow 0$, this implies that the CLT does not hold for the sum.

Chernoff bound

- We have done CLT, but it does not give us explicit tail bounds.
- Lets look at concentration inequalities.

Theorem (Chernoff bound for Bernoullis)

Let $X_i \in \{0, 1\}$ be independent random variables with $E[X_i] = p_i$. Let $X := \sum_i X_i$, $\mu := \sum_i p_i$. For $0 < \delta < 1$,

$$P(X \geq \mu(1 + \delta)) \leq e^{-\delta^2 \mu / 3} \quad P(X \leq \mu(1 - \delta)) \leq e^{-\delta^2 \mu / 2}$$

- How about subgaussian r.v.s?

Sub-Gaussian random variables

Theorem

For X_1, \dots, X_n independent sub-gaussian random variables with sub-gaussian parameters σ_i^2 and $E[X_i] = \mu_i$, for $\forall t > 0$,

$$P\left(\sum_i (X_i - \mu_i) \geq t\right) \leq e^{-\frac{t^2}{2\sum_i \sigma_i^2}}$$

- If $X_i \in [a, b]$, $E[X_i] = 0$, using Hoeffding's lemma we get:

$$P\left(\sum_i X_i \geq t\right) \leq e^{-\frac{2t^2}{n(b-a)^2}}$$

- If $X_i \sim N(0, \sigma^2)$, we immediately get back the chernoff bound for Gaussians.

Sub-exponential random variables

Definition

X is sub-exponential with parameters (ν, b) if, $\forall |\lambda| < 1/b$,

$$\log M_{X-\mu}(\lambda) \leq \frac{\lambda^2 \nu^2}{2}$$

Theorem

Let X be a sub-exponential random variable with parameters (ν, b) . Then,

$$P(X \geq \mu + t) \leq \begin{cases} e^{-\frac{t^2}{2\nu^2}} & \text{if } 0 \leq t \leq \frac{\nu^2}{b} \\ e^{-\frac{t}{2b}} & \text{if } t \geq \frac{\nu^2}{b} \end{cases}$$

- For small t this is sub-gaussian in nature, whereas for large t the exponent decays linearly with t .

Bernstein's condition and the sub-exponential property

Definition

A random variable with mean μ and variance σ^2 satisfies the Bernstein condition with parameter $b > 0$, if $|E[(X - \mu)^k]| \leq \frac{1}{2} k! \sigma^2 b^{k-2}$ for $k \geq 2$.

Theorem

If X ($E[X] = \mu$, $\text{var}(X) = \sigma^2$) satisfies the Bernstein condition with parameter $b > 0$, then X is sub-exponential with $(\sqrt{2}\sigma, 2b)$.

Theorem

If X with mean μ and variance σ^2 satisfies the Bernstein condition with parameter $b > 0$, then

$$P(|X - \mu| \geq t) \leq 2e^{-\frac{t^2}{2(\sigma^2 + bt)}} \quad (5)$$

How about martingale inequalities?

Theorem

Let $f : \mathcal{X}^n \rightarrow \mathcal{R}$ satisfy the following bounded difference condition

$\forall x_1, \dots, x_n, x'_i \in \mathcal{X}$:

$$|f(x_1, \dots, x_{i-1}, x_i, x_{i+1}, \dots, x_n) - f(x_1, \dots, x_{i-1}, x'_i, x_{i+1}, \dots, x_n)| \leq B_i,$$

then, $P(|f(X) - E[f(X)]| \geq t) \leq 2 \exp\left(-\frac{2t^2}{\sum_i B_i^2}\right)$

- Note that this boils down to Hoeffding's when f is the sum of bounded random variables.

Recall-Lipschitz functions of Gaussian random variables

Theorem (LG:Lipschitz functions of Gaussians)

Let (X_1, \dots, X_n) be a vector of iid $N(0, 1)$ random variables. Let $f : \mathcal{R}^n \rightarrow \mathcal{R}$ be L -Lipschitz w.r.t the Euclidean norm. Then $f(X) - E[f(X)]$ is sub-gaussian with parameter at most L , i.e. $\forall t \geq 0$,

$$P(|f(X) - E[f(X)]| \geq t) \leq e^{-\frac{t^2}{2L^2}}$$

- So a L -Lipschitz function of n gaussian random variables behave like a subgaussian with variance proxy L^2 .

Convex Lipschitz functions of bounded random variables

Theorem

Consider a convex function $f : \mathcal{R}^n \rightarrow \mathcal{R}$ with Lipschitz constant L . Also consider n iid random variables $X_1, \dots, X_n \in \{-1, 1\}$. We have for $t > 0$

$$P(|f(X) - M_f| \geq t) \leq 4 \exp \left(-\frac{t^2}{16L^2} \right),$$

where M_f is the median of f .

- Often the median can be replaced by the mean with a little give in the t .

Example

Consider a mean zero iid sequence $X = \{X_i\}_{i=1}^n$. We will bound $f(X) := \sup_{a \in \mathcal{A}} a^T X$ where \mathcal{A} is a compact subset of \mathcal{R}^n such that $\mathcal{W} = \sup_{a \in \mathcal{A}} \|a\|_2 < \infty$.

- If $X_i \sim N(0, 1)$ using Gaussian+Lipschitz
$$P(|f(X) - E[f(X)]| \geq t) \leq 2e^{-\frac{t^2}{2\mathcal{W}^2}}$$
- If X_i are bounded, then Talagrand gives us the same thing (modulo constants).
- How about McDiarmid? Gives a weaker result.
- This is all very good, but how about $E[f(X)]$. If $X \sim N(0, 1)$, this is the Gaussian complexity.

Recall the finite class lemma?

If \mathcal{A} is finite, we can use the following.

Theorem

Consider z with independent sub-gaussian components.

$$E \max_{a \in A} \langle z, a \rangle \leq \max_{a \in A} \|a\| \sqrt{2 \log |A|}$$

- What happens when A is compact, and not finite?
- Use the discretization lemma, or the metric entropy integral!

Upper bound by 1 step discretization

Theorem

(1-step discretization bound). Let $\{X_\theta, \theta \in \mathcal{T}\}$ be a zero-mean sub-Gaussian process with respect to the metric d_X . Then for any $\delta > 0$, we have

$$E \left[\sup_{\theta, \theta' \in \mathcal{T}} (X_\theta - X_{\theta'}) \right] \leq 2E \left[\sup_{\substack{\theta, \theta' \in \mathcal{T} \\ d_X(\theta, \theta') \leq \delta}} (X_\theta - X_{\theta'}) \right] + 2D \sqrt{\log N(\delta; \mathcal{T}, d_X)},$$

where $D := \max_{\theta, \theta' \in \Theta} d_X(\theta, \theta')$.

- The mean zero condition gives us: $E[\sup_{a \in \mathcal{A}} a^T X] \leq E[\sup_{a, a' \in \mathcal{A}} (a^T X - a'^T X)]$
- $a^T X$ is sub Gaussian w.r.t the $\|\cdot\|_2$ norm.
- $D = 2\mathcal{W}$.
- Then optimize. You will also need more information about \mathcal{A} to make sure that you can calculate the covering number.

Putting everything in place

- First we do convergence, since it shows up everywhere.
- Next we look at concentration, for sums of bounded, and unbounded random variables, as long as the tails are well controlled.
- Now you want uniform laws, or uniform error bounds. Why? Say you are looking at convergence of a nonconvex algorithm. You want to understand the behavior of the convergence within some radius of some local/global optima. Here is where uniform error bounds come in very handy.
- In order to do uniform laws, one also needs a handle over the expectations of the supremum. This is why we looked at:
 - Finite class lemma, VC dimension, Sauer's lemma
 - Covering and packing numbers, Chaining, metric entropy.
 - We also saw that covering numbers can be helpful in bounding tails of suprema, not just expectations.
 - As for distributional convergence, we only looked at the Hajek projections, which helped us with U statistics.