

# SDS 385: Stat Models for Big Data

## Lecture 12: PCA and LDA

---

Purnamrita Sarkar  
Department of Statistics and Data Science  
The University of Texas at Austin  
<https://psarkar.github.io/teaching>

# Principal Component Analysis

- Goal: Find the direction of the most variance.
- Say  $X$  is the data matrix
- The average is  $\bar{\mathbf{x}} = \frac{\sum_{i=1}^n \mathbf{x}_i}{n}$
- Let  $\tilde{\mathbf{x}}_i = \mathbf{x}_i - \bar{\mathbf{x}}$


# Principal Component Analysis

- Goal: Find the direction of the most variance.
- Say  $X$  is the data matrix
- The average is  $\bar{\mathbf{x}} = \frac{\sum_{i=1}^n \mathbf{x}_i}{n}$
- Let  $\tilde{\mathbf{x}}_i = \mathbf{x}_i - \bar{\mathbf{x}}$
- The sample variance of  $(\tilde{\mathbf{x}}_1, \dots, \tilde{\mathbf{x}}_n)$  *along a direction*  $w$  is give by:

$$\frac{1}{n} \sum_{i=1}^n (\tilde{\mathbf{x}}_i^T w)^2$$

- What is the sample variance of  $(\mathbf{x}_1, \dots, \mathbf{x}_n)$  *along a direction*  $w$ ?

# Principal Component Analysis



# First component

- So the first PC direction is:

$$\mathbf{w}_1 = \arg \max_{\|\mathbf{w}\|=1} \frac{1}{n} \sum_{i=1}^n (\tilde{\mathbf{x}}_i^T \mathbf{w})^2$$

- And the first PC component of  $\tilde{\mathbf{x}}_i$  is  $\tilde{\mathbf{x}}_i^T \mathbf{w}_1$

# First component

- So the first PC direction is:

$$\mathbf{w}_k = \arg \max_{\substack{\|\mathbf{w}\|=1 \\ \mathbf{w} \perp \mathbf{w}_1, \dots, \mathbf{w}_{k-1}}} \frac{1}{n} \sum_{i=1}^n (\tilde{\mathbf{x}}_i^T \mathbf{w})^2$$

- And the  $k^{th}$  PC component of  $\tilde{\mathbf{x}}_i$  is  $\tilde{\mathbf{x}}_i^T \mathbf{w}_k$
- Note that  $\mathbf{w}_1, \dots, \mathbf{w}_k$  form an orthogonal basis.

# Simple algorithm

- Let  $W$  is a matrix with  $w_k$  along its columns
- $\tilde{X}W$  gives a **low dimensional** representation of  $\tilde{X}$

# Simple algorithm

- Let  $W$  is a matrix with  $\mathbf{w}_k$  along its columns
- $\tilde{X}W$  gives a **low dimensional** representation of  $\tilde{X}$
- We can frame the optimization problem also as

$$\mathbf{w}_1 = \arg \max_{\|\mathbf{w}\|=1} \mathbf{w}^T \tilde{X}^T \tilde{X} \mathbf{w}$$



# Simple algorithm

- Let  $W$  is a matrix with  $\mathbf{w}_k$  along its columns
- $\tilde{X}W$  gives a **low dimensional** representation of  $\tilde{X}$
- We can frame the optimization problem also as

$$\mathbf{w}_1 = \arg \max_{\|\mathbf{w}\|=1} \mathbf{w}^T \tilde{X}^T \tilde{X} \mathbf{w}$$

- This is the first eigenvector of  $S = \tilde{X}^T \tilde{X}$

# Eigenvector and eigenvalues

- Any square symmetric matrix  $S$  has real eigenvalues
- The  $i^{th}$  eigenvalue, vector pair satisfy  $S\mathbf{w}_i = \lambda_i\mathbf{w}_i$
- The eigenvectors are orthogonal to each other, and normalized to have length 1.

# Eigenvector and eigenvalues

- Any square symmetrix matrix  $S$  has real eigenvalues
- The  $i^{th}$  eigenvalue,vector pair satisfy  $S\mathbf{w}_i = \lambda_i\mathbf{w}_i$
- The eigenvectors are orthogonal to each other, and normalized to have length 1.
- In matrix terms, we can write:

$$S = U\Sigma U^T, \text{ where}$$

- columns of  $U$  are the orgonal eigenvectors, and
- $\Sigma$  is a diagonal matrix with eigenvalues on the diagonal
- The larger the magnitude of the eigenvalue, more important the eigenvector

## Back to PCA: Simple algorithm

- Let  $W$  is a matrix with  $w_k$  along its columns
- $\tilde{X}W$  gives a **low dimensional** representation of  $\tilde{X}$

## Back to PCA: Simple algorithm

- Let  $W$  is a matrix with  $\mathbf{w}_k$  along its columns
- $\tilde{X}W$  gives a **low dimensional** representation of  $\tilde{X}$
- We can frame the optimization problem also as

$$\mathbf{w}_1 = \arg \max_{\|\mathbf{w}\|=1} \mathbf{w}^T \tilde{X}^T \tilde{X} \mathbf{w}$$

## Back to PCA: Simple algorithm

- Let  $W$  is a matrix with  $\mathbf{w}_k$  along its columns
- $\tilde{X}W$  gives a **low dimensional** representation of  $\tilde{X}$
- We can frame the optimization problem also as

$$\mathbf{w}_1 = \arg \max_{\|\mathbf{w}\|=1} \mathbf{w}^T \tilde{X}^T \tilde{X} \mathbf{w}$$

- This is the first eigenvector of  $S = \tilde{X}^T \tilde{X}$

## Back to PCA: Simple algorithm

- Let  $W$  is a matrix with  $\mathbf{w}_k$  along its columns
- $\tilde{X}W$  gives a **low dimensional** representation of  $\tilde{X}$
- We can frame the optimization problem also as

$$\mathbf{w}_1 = \arg \max_{\|\mathbf{w}\|=1} \mathbf{w}^T \tilde{X}^T \tilde{X} \mathbf{w}$$

- This is the first eigenvector of  $S = \tilde{X}^T \tilde{X}$
- What is  $S$ ?

## Back to PCA: Simple algorithm

- Let  $W$  is a matrix with  $\mathbf{w}_k$  along its columns
- $\tilde{X}W$  gives a **low dimensional** representation of  $\tilde{X}$
- We can frame the optimization problem also as

$$\mathbf{w}_1 = \arg \max_{\|\mathbf{w}\|=1} \mathbf{w}^T \tilde{X}^T \tilde{X} \mathbf{w}$$

- This is the first eigenvector of  $S = \tilde{X}^T \tilde{X}$
- What is  $S$ ?
- Its the scalar multiple of the sample covariance matrix

$$\hat{\Sigma} = \frac{1}{n} \tilde{\mathbf{x}}_i \tilde{\mathbf{x}}_i^T = \frac{S}{n}$$



## Back to PCA: Simple algorithm

- Let  $W$  is a matrix with  $\mathbf{w}_k$  along its columns
- $\tilde{X}W$  gives a **low dimensional** representation of  $\tilde{X}$
- We can frame the optimization problem also as

$$\mathbf{w}_1 = \arg \max_{\|\mathbf{w}\|=1} \mathbf{w}^T \tilde{X}^T \tilde{X} \mathbf{w}$$

- This is the first eigenvector of  $S = \tilde{X}^T \tilde{X}$
- What is  $S$ ?
- Its the scalar multiple of the sample covariance matrix

$$\hat{\Sigma} = \frac{1}{n} \tilde{\mathbf{x}}_i \tilde{\mathbf{x}}_i^T = \frac{S}{n}$$

- So, all you have to do is to calculate eigenvectors of the covariance matrix.

## Back to PCA: Simple algorithm

- So, all you have to do is to calculate eigenvectors of the covariance matrix.
- But, do I even need to do that?

## Back to PCA: Simple algorithm

- So, all you have to do is to calculate eigenvectors of the covariance matrix.
- But, do I even need to do that?
- The right singular vectors of  $\tilde{X}$  is just fine.

## Back to PCA: Simple algorithm

- So, all you have to do is to calculate eigenvectors of the covariance matrix.
- But, do I even need to do that?
- The right singular vectors of  $\tilde{X}$  is just fine.
- How many PC's? (more of a dissertaiton question)

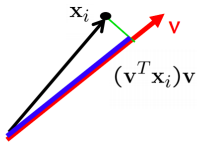
# Singular value decomposition

$$A = U \Sigma V^T$$

$m \times m$     $m \times m$     $V$  is  $m \times n$

- The columns of  $U$  are orthogonal eigenvectors of  $AA^T$
- The columns of  $V$  are orthogonal eigenvectors of  $A^T A$
- $A^T A$  and  $AA^T$  have the same eigenvalues

## Second interpretation



- Minimum reconstruction error:

$$(\mathbf{x}_i - (\mathbf{x}_i^T \mathbf{w}) \mathbf{w})^T (\mathbf{x}_i - (\mathbf{x}_i^T \mathbf{w}) \mathbf{w}) = \mathbf{x}_i^T \mathbf{x}_i - (\mathbf{x}_i^T \mathbf{w})^2$$

- So, the first PC direction gives the direction projecting on which has the **minimum reconstruction error**.

# Low rank approximation

- Take the centered data matrix  $\tilde{X}$  with SVD

$$\tilde{X} = USV^T$$

- Project on the top  $k$  PC's  $W \in \mathbb{R}^{p \times k}$

# Low rank approximation

- Take the centered data matrix  $\tilde{X}$  with SVD

$$\tilde{X} = USV^T$$

- Project on the top  $k$  PC's  $W \in \mathbb{R}^{p \times k}$
- You get  $\tilde{X}W = US_kV^T$ , where  $S_k$  has zeroed out all singular values  $< \sigma_k$



# Low rank approximation

- Take the centered data matrix  $\tilde{X}$  with SVD

$$\tilde{X} = USV^T$$

- Project on the top  $k$  PC's  $W \in \mathbb{R}^{p \times k}$
- You get  $\tilde{X}W = US_kV^T$ , where  $S_k$  has zeroed out all singular values  $< \sigma_k$
- So  $W = \arg \min_{\text{rank}(B)=k, B \in \mathbb{R}^{n \times p}} \|\tilde{X} - B\|_F^2$  and the reconstruction

error is  $\sum_{i=k+1}^p \sigma_i^2$

# Low rank approximation

- Take the centered data matrix  $\tilde{X}$  with SVD

$$\tilde{X} = USV^T$$

- Project on the top  $k$  PC's  $W \in \mathbb{R}^{p \times k}$
- You get  $\tilde{X}W = US_kV^T$ , where  $S_k$  has zeroed out all singular values  $< \sigma_k$

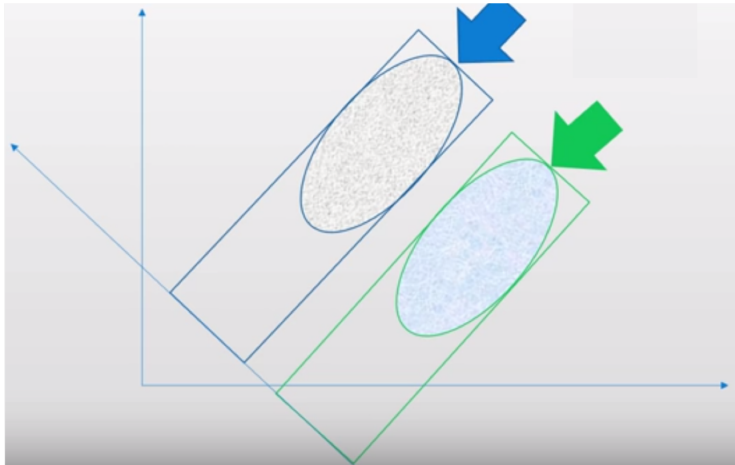
- So  $W = \arg \min_{\text{rank}(B)=k, B \in \mathbb{R}^{n \times p}} \|\tilde{X} - B\|_F^2$  and the reconstruction

error is  $\sum_{i=k+1}^p \sigma_i^2$

- This explains why you want to take large  $k$  to reduce approx. error.

# Linear Discriminant Analysis

- PCA did not have class information
- LDA does take that into account.
- We will do it for two classes.

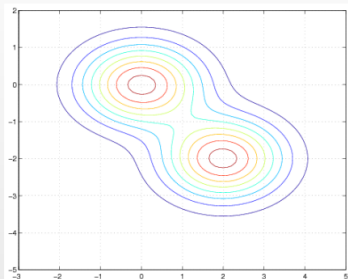
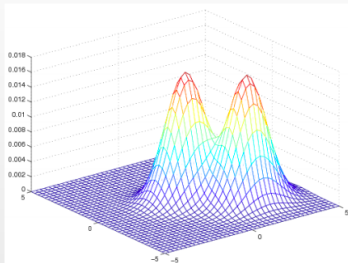


- Assume that the data is coming from a mixture of two Gaussians with parameters  $(\mu_k, \Sigma_k), k \in \{1, 2\}$

- Assume that the data is coming from a mixture of two Gaussians with parameters  $(\mu_k, \Sigma_k), k \in \{1, 2\}$
- Recall the density of a multivariate gaussian

$$f(x) = \frac{1}{(2\pi)^{p/2} |\Sigma_k|^{1/2}} \exp \left( -\frac{1}{2} (x - \mu_k)^T \Sigma_k^{-1} (x - \mu_k) \right)$$

# A pretty picture



- Assign point  $x$  to the class with maximizes the posterior probability of belonging to that class

- Assign point  $\mathbf{x}$  to the class with maximizes the posterior probability of belonging to that class

$$\begin{aligned}\arg \max_k P(y = k | \mathbf{x}, \Theta) &= \arg \max_k \frac{P(\mathbf{x} | y = k, \Theta) P(y = k)}{P(\mathbf{x})} \\ &= \arg \max_k P(\mathbf{x} | y = k, \Theta) P(y = k)\end{aligned}$$



- Assign point  $\mathbf{x}$  to the class with maximizes the posterior probability of belonging to that class

$$\begin{aligned}\arg \max_k P(y = k | \mathbf{x}, \Theta) &= \arg \max_k \frac{P(\mathbf{x} | y = k, \Theta) P(y = k)}{P(\mathbf{x})} \\ &= \arg \max_k P(\mathbf{x} | y = k, \Theta) P(y = k)\end{aligned}$$

- So decision rule for class 1 is

$$\begin{aligned}& -\frac{1}{2} \log |\Sigma_1|^{1/2} - \frac{1}{2} (x - \mu_1)^T \Sigma_1^{-1} (x - \mu_1) + \log \pi_1 \\ & > -\frac{1}{2} \log |\Sigma_2|^{1/2} - \frac{1}{2} (x - \mu_2)^T \Sigma_2^{-1} (x - \mu_2) + \log \pi_2\end{aligned}$$

- For  $\Sigma_1 \neq \Sigma_2$ , this is a quadratic function.

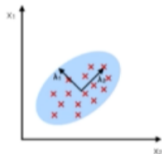
# LDA

- LDA assumes that  $\Sigma_1 = \Sigma_2$
- So now we get a linear decision boundary

$$x^T \Sigma^{-1}(\mu_1 - \mu_2) > \frac{\mu_1 + \mu_2}{2} \Sigma^{-1}(\mu_1 - \mu_2) - \log \frac{\pi_1}{\pi_2}$$

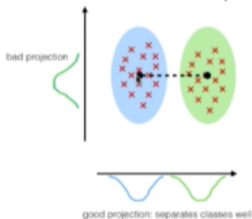
## PCA:

component axes that maximize the variance



## LDA:

maximizing the component axes for class-separation



- Class proportion

$$\hat{\pi}_k = \frac{\sum_{y_i=k} y_i}{n}.$$

- Class proportion

$$\hat{\pi}_k = \frac{\sum_{y_i=k} y_i}{n}.$$

- Class mean

$$\hat{\mu}_k = \frac{\sum_{y_i=k} x_i}{n}.$$

- Class proportion

$$\hat{\pi}_k = \frac{\sum_{y_i=k} y_i}{n}.$$

- Class mean

$$\hat{\mu}_k = \frac{\sum_{y_i=k} x_i}{n}.$$

- Common class covariance matrix

$$\hat{\Sigma} = \frac{\sum_{k=1}^K \sum_{y_i=k} (x_i - \hat{\mu}_k)(x_i - \hat{\mu}_k)^T}{n - K}.$$

- For datapoint  $x$  whose class you want to predict, for each class  $k \in \{1, \dots, K\}$ , compute the **linear discriminant function**  
 $\delta_k(x) = x^T \Sigma^{-1} \mu_k - \frac{1}{2} \mu_k^T \Sigma^{-1} \mu_k + \log \pi_k$  *with estimated parameters*

# Multiple classes

- For datapoint  $x$  whose class you want to predict, for each class  $k \in \{1, \dots, K\}$ , compute the **linear discriminant function**  
$$\delta_k(x) = x^T \Sigma^{-1} \mu_k - \frac{1}{2} \mu_k^T \Sigma^{-1} \mu_k + \log \pi_k$$
 *with estimated parameters*
- Assign  $x$  to class that maximizes this function.

# Multiple classes

- For datapoint  $x$  whose class you want to predict, for each class  $k \in \{1, \dots, K\}$ , compute the **linear discriminant function**  
$$\delta_k(x) = x^T \Sigma^{-1} \mu_k - \frac{1}{2} \mu_k^T \Sigma^{-1} \mu_k + \log \pi_k$$
 *with estimated parameters*
- Assign  $x$  to class that maximizes this function.
- Why not use QDA?



## Multiple classes

- Multiclass LDA minimizes  $\frac{(x - \hat{\mu}_j)^T \hat{\Sigma}^{-1} (x - \hat{\mu}_j)}{2} - \log \hat{\pi}_j$

# Multiple classes

- Multiclass LDA minimizes  $\frac{(x - \hat{\mu}_j)^T \hat{\Sigma}^{-1} (x - \hat{\mu}_j)}{2} - \log \hat{\pi}_j$
- Calculate the square root of  $\hat{\Sigma}^{-1/2}$ 
  - Calculate the eigen-decomposition of  $\hat{\Sigma} = UDU^T$ ,
  - $\hat{\Sigma}^{-1/2} = UD^{-1/2}$

## Multiple classes

- Multiclass LDA minimizes  $\frac{(x - \hat{\mu}_j)^T \hat{\Sigma}^{-1} (x - \hat{\mu}_j)}{2} - \log \hat{\pi}_j$
- Calculate the square root of  $\hat{\Sigma}^{-1/2}$ 
  - Calculate the eigen-decomposition of  $\hat{\Sigma} = UDU^T$ ,
  - $\hat{\Sigma}^{-1/2} = UD^{-1/2}$
- Transform datapoint  $x$  to  $\tilde{x} := D^{-1/2}U^T x$
- Mean becomes:  $\tilde{\mu}_j := D^{-1/2}U^T \hat{\mu}_j$

# Multiple classes

- Multiclass LDA minimizes  $\frac{(x - \hat{\mu}_j)^T \hat{\Sigma}^{-1} (x - \hat{\mu}_j)}{2} - \log \hat{\pi}_j$
- Calculate the square root of  $\hat{\Sigma}^{-1/2}$ 
  - Calculate the eigen-decomposition of  $\hat{\Sigma} = UDU^T$ ,
  - $\hat{\Sigma}^{-1/2} = UD^{-1/2}$
- Transform datapoint  $x$  to  $\tilde{x} := D^{-1/2}U^T x$
- Mean becomes:  $\tilde{\mu}_j := D^{-1/2}U^T \hat{\mu}_j$
- Remember Mahanobis distance?

$$(x - \hat{\mu}_j)^T \hat{\Sigma}^{-1} (x - \hat{\mu}_j) = (\tilde{x} - \tilde{\mu}_j)^T (\tilde{x} - \tilde{\mu}_j)$$

# Multiple classes

- Multiclass LDA minimizes  $\frac{(x - \hat{\mu}_j)^T \hat{\Sigma}^{-1} (x - \hat{\mu}_j)}{2} - \log \hat{\pi}_j$
- Calculate the square root of  $\hat{\Sigma}^{-1/2}$ 
  - Calculate the eigen-decomposition of  $\hat{\Sigma} = UDU^T$ ,
  - $\hat{\Sigma}^{-1/2} = UD^{-1/2}$
- Transform datapoint  $x$  to  $\tilde{x} := D^{-1/2}U^T x$
- Mean becomes:  $\tilde{\mu}_j := D^{-1/2}U^T \hat{\mu}_j$
- Remember Mahanobis distance?

$$(x - \hat{\mu}_j)^T \hat{\Sigma}^{-1} (x - \hat{\mu}_j) = (\tilde{x} - \tilde{\mu}_j)^T (\tilde{x} - \tilde{\mu}_j)$$

- After this “whitening”, the decision rule becomes very simple:
  - Assign  $x$  to class  $j$  such that  $(\tilde{x} - \tilde{\mu}_j)^T (\tilde{x} - \tilde{\mu}_j) - \log \hat{\pi}_j$

# LDA algorithm

- Estimate parameters by  $\hat{\pi}_j, \hat{\mu}_j, \hat{\Sigma}$
- Compute eigendecomposition of  $\hat{\Sigma} = UDU^T$
- Transform the means to  $\tilde{\mu}_j$
- For a datapoint  $x$ , compute the whitened point  $\tilde{x}$
- Now assign to class  $j$  that minimizes  $\frac{1}{2}\text{dist}(\tilde{x}, \tilde{\mu}_j)^2 - \log \hat{\pi}_j$

- How many dimensions do we need to represent 2 points?

- How many dimensions do we need to represent 2 points?
  - Just 1



# LDA and dimensionality reduction

- How many dimensions do we need to represent 2 points?
  - Just 1
- How many dimensions do we need to represent  $K$  means?

# LDA and dimensionality reduction

- How many dimensions do we need to represent 2 points?
  - Just 1
- How many dimensions do we need to represent  $K$  means?
  - Just  $K - 1$

## LDA and dimensionality reduction

- Whiten the data
- Create the matrix  $M$  of means  $[\tilde{\mu}_1 \dots \tilde{\mu}_K]$

# LDA and dimensionality reduction

- Whiten the data
- Create the matrix  $M$  of means  $[\tilde{\mu}_1 \dots \tilde{\mu}_K]$
- Do a PCA of this matrix, and call the top  $K$  singular vectors  $A \in \mathbb{R}^{p \times K}$  and all the rest as  $A_{\perp}$ .

# LDA and dimensionality reduction

- Whiten the data
- Create the matrix  $M$  of means  $[\tilde{\mu}_1 \dots \tilde{\mu}_K]$
- Do a PCA of this matrix, and call the top  $K$  singular vectors  $A \in \mathbb{R}^{p \times K}$  and all the rest as  $A_{\perp}$ .
- For  $\tilde{x}$ , compute  $Ax$

$$\begin{aligned}\|\tilde{x} - \tilde{\mu}_j\|^2 &= \|AA^T\tilde{x} + A_{\perp}A_{\perp}^T\tilde{x} - \tilde{\mu}_j\|^2 \\ &= \|AA^T\tilde{x} - \tilde{\mu}_j\|^2 + \underbrace{\|A_{\perp}A_{\perp}^T\tilde{x}\|^2}_{\text{Does not depend on } j}\end{aligned}$$

# LDA and dimensionality reduction

- Whiten the data
- Create the matrix  $M$  of means  $[\tilde{\mu}_1 \dots \tilde{\mu}_K]$
- Do a PCA of this matrix, and call the top  $K$  singular vectors  $A \in \mathbb{R}^{p \times K}$  and all the rest as  $A_{\perp}$ .
- For  $\tilde{x}$ , compute  $Ax$

$$\begin{aligned}\|\tilde{x} - \tilde{\mu}_j\|^2 &= \|AA^T\tilde{x} + A_{\perp}A_{\perp}^T\tilde{x} - \tilde{\mu}_j\|^2 \\ &= \|AA^T\tilde{x} - \tilde{\mu}_j\|^2 + \underbrace{\|A_{\perp}A_{\perp}^T\tilde{x}\|^2}_{\text{Does not depend on } j}\end{aligned}$$

- So the LDA decision rule will be unchanged if we project into the subspace spanned by the centers.

# Acknowledgment

- Some pictures are borrowed from Brett Bernstein's notes from NYU and Jia Li's notes from PSU
- Some slides are borrowed from Ryan Tibshirani's notes
- Elements of statistical learning, HTF