

# ECS289: Scalable Machine Learning

Cho-Jui Hsieh  
UC Davis

Oct 15, 2015








# Outline

- Matrix Completion (Background)
- Alternating Least Squares (ALS)
- Stochastic Gradient method (SG)
- Coordinate Descent (CD)

# Recommender Systems

## Rating Matrix

Users

	Movie 1	Movie 2	Items							Movie 10	Movie 11
 Hitang-Fu	1			5			3		5		2
 Chao-Jun		2		3			5		2	5	
 Si Si					3	?	5		3		
 Indrajit	2		5			3		4		2	
 Kai-Yang				5			5				1
 Donghyuk		5			1				5		
 Naga	1			1				2			4

# Matrix Factorization Approach $A \approx WH^T$

$H^T$

-0.07	-0.11	-0.53	-0.46	-0.06	-0.05	-0.53	-0.07	-0.35	-0.19	-0.14
0.13	-0.42	0.45	0.17	-0.25	-0.17	-0.18	0.27	-0.59	0.05	0.14
-0.21	-0.43	-0.23	0.16	0.08	0.17	0.57	-0.39	-0.37	-0.08	-0.15

$W$

-8.72	0.03	-1.03
-7.56	-0.79	0.62
-4.07	-3.95	2.55
-3.52	3.73	-3.32
-7.78	2.34	2.33
-2.44	-5.29	-3.92
-1.78	1.90	-1.68

1			5			3		5		2
	2		3			5		2	5	
				3		5		3		
2		5			3		4		2	
			5			5				1
	5			1				5		
1			1				2			4

# Matrix Factorization Approach $A \approx WH^T$

$H^T$

-0.07	-0.11	-0.53	-0.46	-0.06	-0.05	-0.53	-0.07	-0.35	-0.19	-0.14
0.13	-0.42	0.45	0.17	-0.25	-0.17	-0.18	0.27	-0.59	0.05	0.14
-0.21	-0.43	-0.23	0.16	0.08	0.17	0.57	-0.39	-0.37	-0.08	-0.15

$W$

-8.72	0.03	-1.03
-7.56	-0.79	0.62
-4.07	-3.95	2.55
-3.52	3.73	-3.32
-7.78	2.34	2.33
-2.44	-5.29	-3.92
-1.78	1.90	-1.68

1			5			3		5		2
	2		3			5		2	5	
				3	?	5		3		
2		5			3		4		2	
			5			5				1
	5			1				5		
1			1				2			4

# Matrix Factorization Approach


$$\min_{\substack{W \in \mathbb{R}^{m \times k} \\ H \in \mathbb{R}^{n \times k}}} \sum_{(i,j) \in \Omega} (A_{ij} - \mathbf{w}_i^T \mathbf{h}_j)^2 + \lambda (\|W\|_F^2 + \|H\|_F^2),$$

- $\Omega = \{(i, j) \mid A_{ij} \text{ is observed}\}$
- Regularized terms to avoid over-fitting


Matrix factorization maps users/items to **latent feature space**  $\mathbb{R}^k$

- the  $i^{\text{th}}$  user  $\Rightarrow i^{\text{th}}$  row of  $W$ ,  $\mathbf{w}_i^T$ ,
- the  $j^{\text{th}}$  item  $\Rightarrow j^{\text{th}}$  row of  $H$ ,  $\mathbf{h}_j^T$ .
- $\mathbf{w}_i^T \mathbf{h}_j$ : measures the interaction between  $i^{\text{th}}$  user and  $j^{\text{th}}$  item.

# Latent Feature Space



# Latent Feature Space







# Other Factorizations

## Nonnegative Matrix Factorization

$$\min_{W \geq 0, H \geq 0} \|A - WH^T\|_F^2 + \lambda \|W\|_F^2 + \lambda \|H\|_F^2$$

- Each entry is positive
- $A$  is either fully or partially observed
- Goal: find nonnegative latent factors


# NMF vs PCA



# Optimization for Matrix Completion: Alternating Least Squares

# Properties of the Objective Function

- Nonconvex problem (why?)
- Example:  $f(x, y) = \frac{1}{2}(xy - 1)^2$   
 $\nabla f(0, 0) = \mathbf{0}$ , but clearly  $[0, 0]$  is not a global optimum



# ALS: Alternating Least Squares

- Objective function:

$$\min_{W, H} \left\{ \frac{1}{2} \sum_{i,j \in \Omega} (A_{ij} - (WH^T)_{ij})^2 + \frac{\lambda}{2} \|W\|_F^2 + \frac{\lambda}{2} \|H\|_F^2 \right\} := f(W, H)$$

- Iteratively fix either  $H$  or  $W$  and optimize the other:

Input: partially observed matrix  $A$ , initial values of  $W, H$

For  $t = 1, 2, \dots$

Fix  $W$  and update  $H$ :  $H \leftarrow \operatorname{argmin}_H f(W, H)$

Fix  $H$  and update  $W$ :  $W \leftarrow \operatorname{argmin}_W f(W, H)$

# ALS: Alternating Least Squares

- Define:  $\Omega_j := \{i \mid (i, j) \in \Omega\}$
- $\mathbf{w}_i$ : the  $i$ -th row of  $W$ ;  $\mathbf{h}_j$ : the  $j$ -th row of  $H$
- The subproblem:

$$\begin{aligned} \operatorname{argmin}_{\mathbf{H}} \frac{1}{2} \sum_{i,j \in \Omega} (A_{ij} - (WH^T)_{ij})^2 + \frac{\lambda}{2} \|\mathbf{H}\|_F^2 \\ = \sum_{j=1}^n \left( \underbrace{\frac{1}{2} \sum_{i \in \Omega_j} (A_{ij} - \mathbf{w}_i^T \mathbf{h}_j)^2 + \frac{\lambda}{2} \|\mathbf{h}_j\|^2}_{\text{ridge regression problem}} \right) \end{aligned}$$

# ALS: Alternating Least Squares

- Define:  $\Omega_j := \{i \mid (i, j) \in \Omega\}$
- $\mathbf{w}_i$ : the  $i$ -th row of  $W$ ;  $\mathbf{h}_j$ : the  $j$ -th row of  $H$
- The subproblem:

$$\begin{aligned} \underset{H}{\operatorname{argmin}} \quad & \frac{1}{2} \sum_{i,j \in \Omega} (A_{ij} - (WH^T)_{ij})^2 + \frac{\lambda}{2} \|H\|_F^2 \\ &= \sum_{j=1}^n \underbrace{\left( \frac{1}{2} \sum_{i \in \Omega_j} (A_{ij} - \mathbf{w}_i^T \mathbf{h}_j)^2 + \frac{\lambda}{2} \|\mathbf{h}_j\|^2 \right)}_{\text{ridge regression problem}} \end{aligned}$$

- $n$  ridge regression problems, each with  $k$  variables  
 $\Rightarrow O(|\Omega|k^2 + nk^3)$
- Easy to parallelize ( $n$  independent ridge regression subproblems)

# ALS: Alternating Least Squares

$$\begin{pmatrix} H^T \end{pmatrix} \begin{pmatrix} w_1^T \\ w_2^T \\ w_3^T \end{pmatrix} \begin{pmatrix} A_{11} & A_{12} & A_{13} \\ A_{21} & A_{22} & A_{23} \\ A_{31} & A_{32} & A_{33} \end{pmatrix}$$
$$\begin{pmatrix} H^T \end{pmatrix} \begin{pmatrix} w_1^T \\ w_2^T \\ w_3^T \end{pmatrix} \begin{pmatrix} A_{11} & A_{12} & A_{13} \\ A_{21} & A_{22} & A_{23} \\ A_{31} & A_{32} & A_{33} \end{pmatrix}$$



# Optimization for Matrix Completion: Stochastic Gradient Method

# Stochastic Gradient Method

- $n_i^W$  : number of nonzeros in the  $i$ -th row of  $A$   
 $n_j^H$  : number of nonzeros in the  $j$ -th column of  $A$
- Decompose the problem into  $\Omega$  components:

$$\begin{aligned} f(W, H) &= \frac{1}{2} \sum_{i,j \in \Omega} (A_{ij} - \mathbf{w}_i^T \mathbf{h}_j)^2 + \frac{\lambda}{2} \|W\|_F^2 + \frac{\lambda}{2} \|H\|_F^2 \\ &= \frac{1}{|\Omega|} \sum_{i,j \in \Omega} \underbrace{\left( \frac{|\Omega|}{2} (A_{ij} - \mathbf{w}_i^T \mathbf{h}_j)^2 + \frac{\lambda |\Omega|}{2 n_i^W} \|\mathbf{w}_i\|^2 + \frac{\lambda |\Omega|}{2 n_j^H} \|\mathbf{h}_j\|^2 \right)}_{f_{i,j}(W, H)} \end{aligned}$$

# Stochastic Gradient Method

- $n_i^W$  : number of nonzeros in the  $i$ -th row of  $A$   
 $n_j^H$  : number of nonzeros in the  $j$ -th column of  $A$
- Decompose the problem into  $\Omega$  components:

$$\begin{aligned} f(W, H) &= \frac{1}{2} \sum_{i,j \in \Omega} (A_{ij} - \mathbf{w}_i^T \mathbf{h}_j)^2 + \frac{\lambda}{2} \|W\|_F^2 + \frac{\lambda}{2} \|H\|_F^2 \\ &= \frac{1}{|\Omega|} \sum_{i,j \in \Omega} \underbrace{\left( \frac{|\Omega|}{2} (A_{ij} - \mathbf{w}_i^T \mathbf{h}_j)^2 + \frac{\lambda |\Omega|}{2 n_i^W} \|\mathbf{w}_i\|^2 + \frac{\lambda |\Omega|}{2 n_j^H} \|\mathbf{h}_j\|^2 \right)}_{f_{i,j}(W, H)} \end{aligned}$$

- The gradient of each component:

$$\begin{aligned} \nabla_{\mathbf{w}_i} f_{i,j}(W, H) &= |\Omega| (\mathbf{w}_i^T \mathbf{h}_j - A_{ij}) \mathbf{h}_j + \frac{\lambda |\Omega|}{n_i^W} \mathbf{w}_i \\ \nabla_{\mathbf{h}_j} f_{i,j}(W, H) &= |\Omega| (\mathbf{w}_i^T \mathbf{h}_j - A_{ij}) \mathbf{w}_i + \frac{\lambda |\Omega|}{n_j^H} \mathbf{h}_j \end{aligned}$$

# Stochastic Gradient Method

- SG algorithm:

Input; partially observed matrix  $A$ , initial values of  $W, H$

For  $t = 1, 2, \dots$

Randomly pick a pair  $(i, j) \in \Omega$

$$\mathbf{w}_i \leftarrow (1 - \frac{\eta_t \lambda}{n_i^W}) \mathbf{w}_i - \eta_t (\mathbf{w}_i^T \mathbf{h}_j - A_{ij}) \mathbf{h}_j$$

$$\mathbf{h}_j \leftarrow (1 - \frac{\eta_t \lambda}{n_j^H}) \mathbf{h}_j - \eta_t (\mathbf{w}_i^T \mathbf{h}_j - A_{ij}) \mathbf{w}_i$$

# Stochastic Gradient Method

- SG algorithm:

Input; partially observed matrix  $A$ , initial values of  $W, H$

For  $t = 1, 2, \dots$

Randomly pick a pair  $(i, j) \in \Omega$

$$\mathbf{w}_i \leftarrow (1 - \frac{\eta_t \lambda}{n_i^W}) \mathbf{w}_i - \eta_t (\mathbf{w}_i^T \mathbf{h}_j - A_{ij}) \mathbf{h}_j$$

$$\mathbf{h}_j \leftarrow (1 - \frac{\eta_t \lambda}{n_j^H}) \mathbf{h}_j - \eta_t (\mathbf{w}_i^T \mathbf{h}_j - A_{ij}) \mathbf{w}_i$$

- Time complexity:  $O(k)$  per iteration;  $O(|\Omega|k)$  for one pass of all observed entries.

# Stochastic Gradient Method

$$\begin{pmatrix} h_1 & h_2 & h_3 \end{pmatrix}$$

$$\begin{pmatrix} w_1^T \\ w_2^T \\ w_3^T \end{pmatrix} \begin{pmatrix} A_{11} & A_{12} & A_{13} \\ A_{21} & A_{22} & A_{23} \\ A_{31} & A_{32} & A_{33} \end{pmatrix}$$

$$\begin{pmatrix} h_1 & h_2; & h_3 \end{pmatrix}$$

$$\begin{pmatrix} w_1^T \\ w_2^T \\ w_3^T \end{pmatrix} \begin{pmatrix} A_{11} & A_{12} & A_{13} \\ A_{21} & A_{22} & A_{23} \\ A_{31} & A_{32} & A_{33} \end{pmatrix}$$

# Optimization for Matrix Completion:

## Distributed Stochastic Gradient Descent (DSGD)

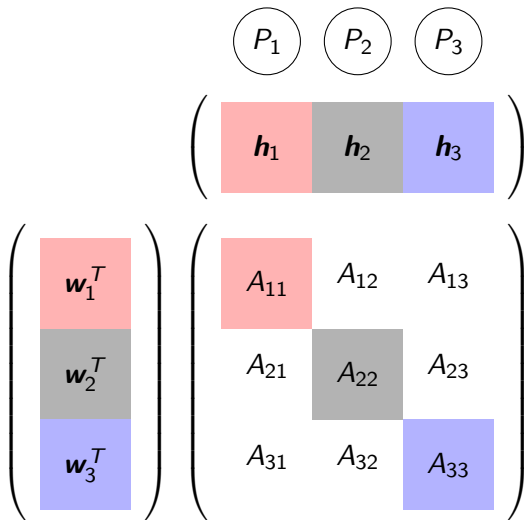
# How to parallelize SG?

- Two SG updates on  $(i_1, j_1)$  and  $(i_2, j_2)$  in the same time:
  - $(i_1, j_1)$ : Update  $\mathbf{w}_{i_1}$  and  $\mathbf{h}_{j_1}$
  - $(i_2, j_2)$ : Update  $\mathbf{w}_{i_2}$  and  $\mathbf{h}_{j_2}$
- Conflict happens when  $i_1 = i_2$  or  $j_1 = j_2$
- How to avoid confliction?

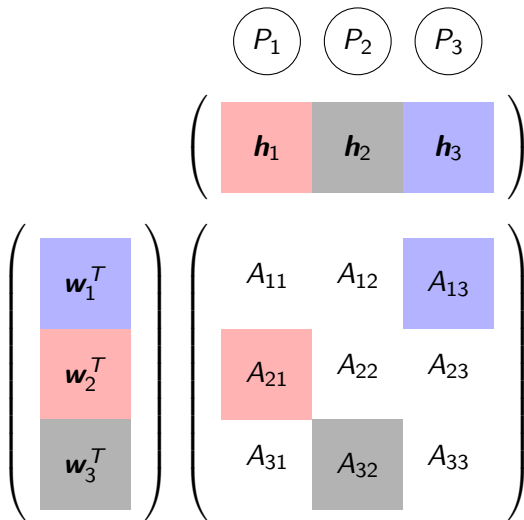
Gemulla et al., “Large-Scale Matrix Factorization with Distributed Stochastic Gradient Descent”. In KDD 2011.



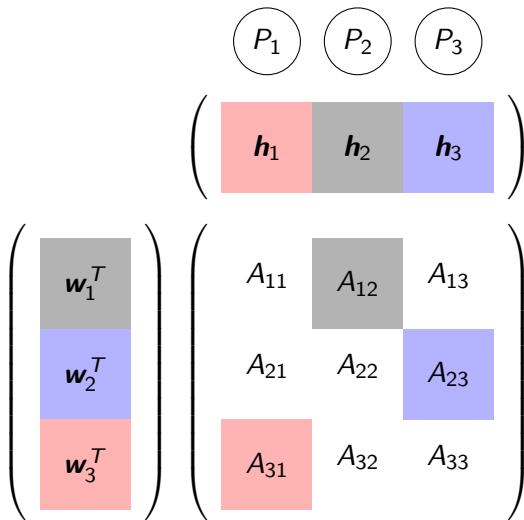
# DSGD: Distributed SGD [Gemulla et al, 2011]



# DSGD: Distributed SGD



# DSGD: Distributed SGD



# Optimization for Matrix Completion: Coordinate Descent

# Coordinate Descent

Update a variable at a time:


$$w_{it} \leftarrow \frac{\sum_{j \in \Omega_i} (A_{ij} - \mathbf{w}_i^T \mathbf{h}_j + w_{it} h_{jt}) h_{jt}}{\lambda + \sum_{j \in \Omega_i} h_{jt}^2}.$$

- Subproblem is just a univariate quadratic problem
- $\Omega_i = \{j : (i, j) \in \Omega\}$
- Can be done in  $O(|\Omega_i|)$


Update Sequence:

- Item/user-wise update:
  - pick a user  $i$  or an item  $j$
  - update the  $i$ -th row of  $W$  or the  $j$ -th column of  $H$
- Feature-wise update:
  - pick a feature index  $t \in \{1, \dots, k\}$
  - update  $t$ -column of  $W$  and  $H$  alternatively


# Feature-wise Update: CCD++



# Feature-wise Update: CCD++




# Feature-wise Update: CCD++







# Feature-wise Update: CCD++




# Feature-wise Update: CCD++




# Feature-wise Update: CCD++




# Feature-wise Update: CCD++




# Feature-wise Update: CCD++




# Feature-wise Update: CCD++



# Feature-wise Update: CCD++




# Feature-wise Update: CCD++






# Feature-wise Update: CCD++




# Feature-wise Update: CCD++



- Cycle through  $k$  feature dimensions

netflix with  $k = 40$



# Coming up

- Next class: other matrix completion topics

Questions?