

SDS 385: Stat Models for Big Data

Lecture 11: Map reduce

Purnamrita Sarkar
Department of Statistics and Data Science
The University of Texas at Austin

<https://psarkar.github.io/teaching>

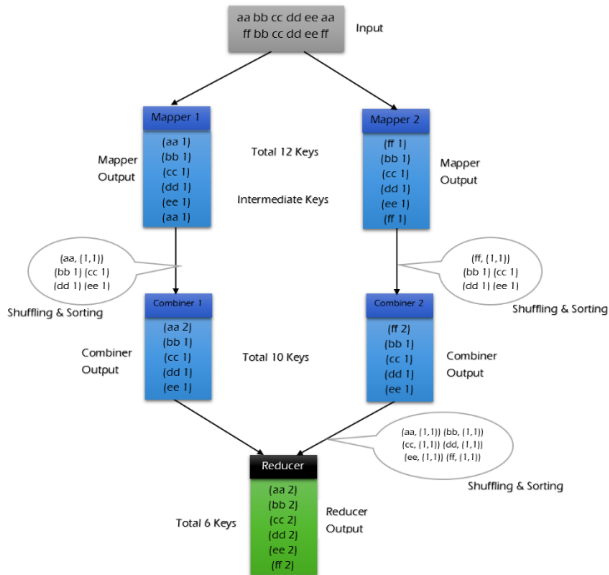
Map reduce framework

- Provides an easy framework for parallel computation
 - distributes data and takes care of synchronization
- But there are restrictions: only certain kinds of parallelisms are allowed

Map reduce basics

- Represent data as $\langle key, value \rangle$ pairs
 - Typically “key” is simple, e.g. the url of an webpage/ ID of a document
 - “value” is complex, can be contents of a webpage/ a document
- distributes data and takes care of synchronization
 - Data from the mapper is fed to the combiner, which is like a local reducer and reduces the amount of data transfer
- But there are restrictions: only certain kinds of parallelisms are allowed

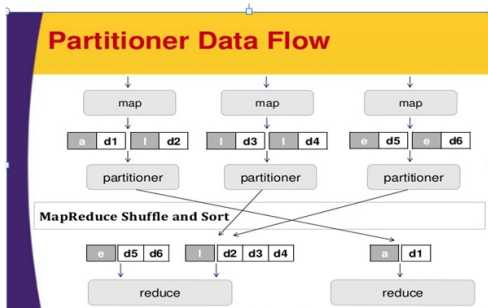
Map reduce basics



Map reduce basics - partitioner

- Runs on the same machine where the mapper had completed its execution by consuming the mapper output.
- Distributes how outputs from the map stage are sent to the reducers.
- Controls the keys partition of the intermediate map-outputs.
- The key is used to derive the partition by a hash function.

Map reduce basics



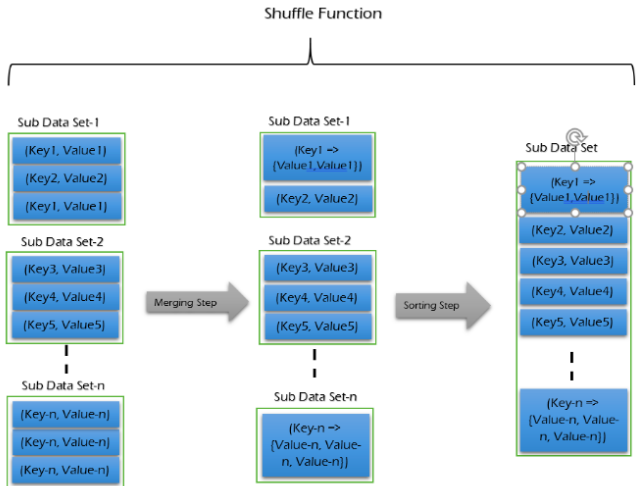
Shuffle and sort

- The shuffling process starts right away as the first mapper has completed its task.
- Once the data is shuffled to the reducer node the intermediate output is sorted based on key before sending it to reduce task.
- The algorithm used for sorting at reducer node is Merge sort.

Shuffle and sort



- The shuffling process starts right away as the first mapper has completed its task.
- Once the data is shuffled to the reducer node the intermediate output is sorted based on key before sending it to reduce task.
- The algorithm used for sorting at reducer node is Merge sort.

Shuffle and sort





Example - calculate degree in weighted graph

- Data: (src,dst, weight)
- Key (src,dst) value weight



Example - calculate degree in weighted graph

- Data: (src,dst, weight)
- Key (src,dst) value weight




Example - transpose of a giant matrix

- Input: sparse matrix in row major order
- Output: sparse matrix in column major order

Map:

- Input: $\langle \text{row } i, (\text{col}_{i1}, \text{val}_{i1}), (\text{col}_{i2}, \text{val}_{i2}), \dots \rangle$
- Output: $\langle \text{col}_{i1}, (\text{row } i, \text{val}_{i1}) \rangle$
- $\langle \text{col}_{i2}, (\text{row } i, \text{val}_{i2}) \rangle$
-

a	b	
	c	d
	e	



Example - transpose of a giant matrix


- Input: sparse matrix in row major order
- Output: sparse matrix in column major order

Map:

- Input: $\langle \text{row } i, (\text{col}_{i1}, \text{val}_{i1}), (\text{col}_{i2}, \text{val}_{i2}), \dots \rangle$
- Output: $\langle \text{col}_{i1}, (\text{row } i, \text{val}_{i1}) \rangle$
- $\langle \text{col}_{i2}, (\text{row } i, \text{val}_{i2}) \rangle$
-

a	b	
	c	d
	e	

Shuffle:



Example - transpose of a giant matrix

- Input: sparse matrix in row major order
- Output: sparse matrix in column major order

Map:


- Input: $\langle \text{row } i, (\text{col}_{i1}, \text{val}_{i1}), (\text{col}_{i2}, \text{val}_{i2}), \dots \rangle$
- Output: $\langle \text{col}_{i1}, (\text{row } i, \text{val}_{i1}) \rangle$
- $\langle \text{col}_{i2}, (\text{row } i, \text{val}_{i2}) \rangle$
-

a	b	
	c	d
	e	

Shuffle

Reduce:

- Sort by row number



Example - linear regression

- Data is $x_i = (x_{i1}, \dots, x_{ip})$
- response variable y_i
- Want to compute $w = (X^T X)^{-1} X^T \mathbf{y}$

Example - linear regression

- Data is $\mathbf{x}_i = (x_{i1}, \dots, x_{ip})$
- response variable y_i
- Want to compute $\mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$
- $\mathbf{X}^T \mathbf{X} = \sum_i \mathbf{x}_i \mathbf{x}_i^T$

Example - linear regression

- Data is $\mathbf{x}_i = (x_{i1}, \dots, x_{ip})$
- response variable y_i
- Want to compute $\mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$
- $\mathbf{X}^T \mathbf{X} = \sum_i \mathbf{x}_i \mathbf{x}_i^T$
- $\mathbf{X}^T \mathbf{Y} = \sum_i y_i \mathbf{x}_i$

Example - linear regression

- Input data $\langle k, \{\mathbf{x}_i, y_i\}_{i \in \text{subgroup } k} \rangle$

Example - linear regression

- Input data $\langle k, \{\mathbf{x}_i, y_i\}_{i \in \text{subgroup } k} \rangle$
- Output: two keys
 - K_1 - for $X^T X$ (calculated over subgroup), Value is a $p + 1 \times p + 1$ matrix
 - K_2 - for $X^T I; y$ (calculated over subgroup), Value is a $p + 1 \times 1$ vector

Example - linear regression

- Input data $\langle k, \{\mathbf{x}_i, y_i\}_{i \in \text{subgroup } k} \rangle$
- Output: two keys
 - K_1 - for $X^T X$ (calculated over subgroup), Value is a $p + 1 \times p + 1$ matrix
 - K_2 - for $X^T \mathbf{y}$ (calculated over subgroup), Value is a $p + 1 \times 1$ vector
- Reducer
 - Sum the (K_1, Value) pairs to get $X^T X$
 - Sum the (K_2, Value) pairs to get $X^T \mathbf{y}$
 - Calculate $w = (X^T X)^{-1} X^T \mathbf{y}$

Example - Naive Bayes

- Data $x = (x_1, \dots, x_p)$, discrete
- Response variable y_i , in $1, \dots, K$
- Prediction for datapoint (a_1, \dots, a_p) :
$$y^* = \arg \max_k P(y = k) \prod_j P(x_j = a_j | y = k)$$

Example - Naive Bayes

- Map:

- Input: $\langle k, \{\mathbf{x}_i, y_i\}_{i \in \text{subgroup } k} \rangle$

- Output: three types of keys

- $K_1 = (j, a_j, k)$, and Value1 $\sum_{\text{subgroup}} 1(x_j = a_j, y = k)$

- $K_2 = (k)$ and Value2 $\sum_{\text{subgroup}} 1(y = k)$

- K_3 and Value3 $\sum_{\text{subgroup}} 1$

Example - Naive Bayes

- Reduce

- Estimate $P(x_j = a_j | y = k)$ by $\frac{\sum_{K_1=(j, a_j, k)} \text{Value1}}{\sum_{K_2=k} \text{Value2}}$

- Estimate $P(y = k)$ by $\frac{\sum_{K_2=k} \text{Value2}}{\sum_{K_3} \text{Value3}}$

Example - K-means

- Data $x = (x_1, \dots, x_p)$
- Goal: find clustering that minimizes sum of squares error.
- At any step, we do:
 - Update $y_i^{(t+1)} \leftarrow$ label of the closest center.
 - Update $\mu_k^{(t+1)} \leftarrow \frac{\sum_{i: y_i^{(t+1)}=k} x_i}{\sum_{i: y_i^{(t)}=k} 1}$

Example - K-means

- Map:
 - Input x
 - $k = \operatorname{argmin}_j \|x - \mu_j^{(t)}\|$
 - Output $(k, x, 1)$

Example - K-means

- Map:

- Input x
- $k = \operatorname{argmin}_j \|x - \mu_j^{(t)}\|$
- Output $(k, x, 1)$

- Reduce:

- Reduce into k values of the form: $(j, \underbrace{\sum_{key=j} value_1}_{N_j}, \underbrace{\sum_{key=j} value_2}_{D_j})$
- Estimate $\mu_k^{(t+1)}$ by $\frac{N_k}{D_k}$

Seems like we can do anything

- Any ML algorithm that has an inner step with sum over matrices or vectors concerning n datapoints can be parallelized.

Seems like we can do anything

- Any ML algorithm that has an inner step with sum over matrices or vectors concerning n datapoints can be parallelized.
- Works well when p is small, since then the communication overhead is relatively small
- Since we are invoking MR framework over each iteration of an algorithm (like GD), for large p things are pretty bad.
- So, how to reduce communication?

Acknowledgment

- <https://www.tutorialscampus.com/tutorials/map-reduce/shuffle-and-sort.htm>
- Sergei Vassilvitskii's notes
- Modeling with Hadoop - KDD 2011 tutorial
- Map-Reduce for Machine Learning on Multicore - NIPS 2007 paper by Chu et al.